

Processes

Processes

- ▶ Process Concept
- ▶ Process Scheduling
- ▶ Operations on Processes
- ▶ Interprocess Communication
- ▶ Examples of IPC Systems
- ▶ Communication in Client-Server Systems

Outcomes

- ▶ Learn the basic concepts of Process and its components.
- ▶ State the Process life cycle.

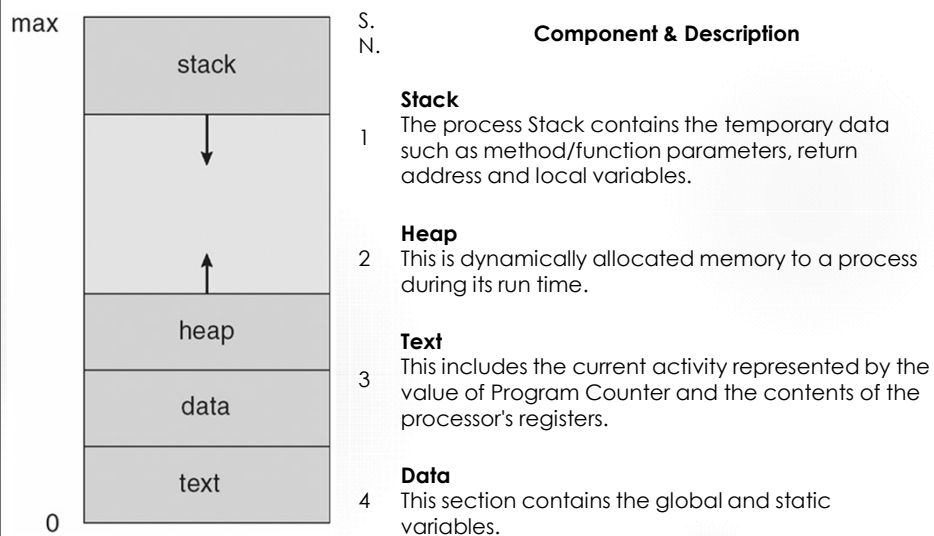
Process Concept

- ▶ An operating system executes a variety of programs:
 - ▶ Batch system – **jobs**
 - ▶ Time-shared systems – **user programs** or **tasks**
- ▶ **Process** – a program in execution; process execution must progress in sequential fashion
- ▶ Multiple parts
 - ▶ The program code, also called **text section**
 - ▶ Current activity including **program counter**, processor registers
 - ▶ **Stack** containing temporary data
 - ▶ Function parameters, return addresses, local variables
 - ▶ **Data section** containing global variables
 - ▶ **Heap** containing memory dynamically allocated during run time

Process Concept (Cont.)

- ▶ A process is basically a program in execution. The execution of a process must progress in a sequential fashion.
- ▶ A process is defined as an entity which represents the basic unit of work to be implemented in the system
- ▶ Program is **passive** entity stored on disk (**executable file**), process is **active**
 - ▶ Program becomes process when executable file loaded into memory
- ▶ Execution of program started via GUI mouse clicks, command line entry of its name, etc
- ▶ One program can be several processes
 - ▶ Consider multiple users executing the same program

Process in Memory



Program

A program is a piece of code which may be a single line or millions of lines. A computer program is usually written by a computer programmer in a programming language. For example, here is a simple program written in C programming language –

```
#include <stdio.h>

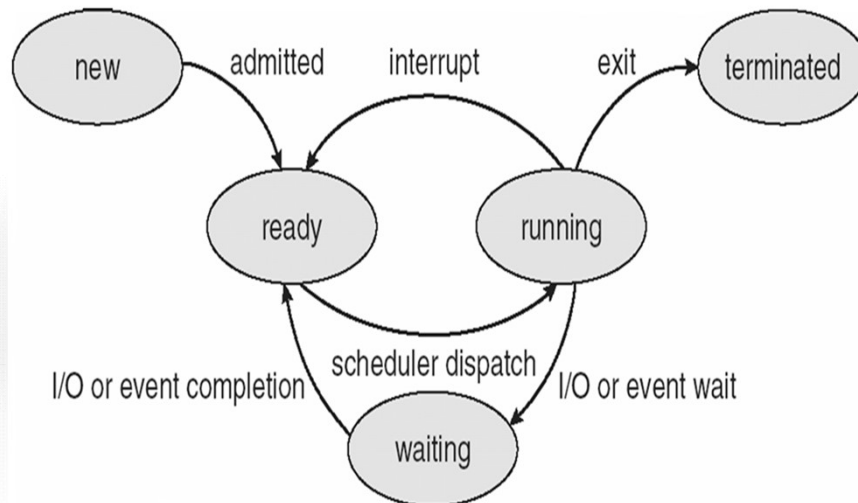
int main() {
    printf("Hello, World! \n");
    return 0;
}
```

A computer program is a collection of instructions that performs a specific task when executed by a computer. When we compare a program with a process, we can conclude that a process is a dynamic instance of a computer program.

Process States

- 1 **Start**
This is the initial state when a process is first started/created.
- 2 **Ready**
The process is waiting to be assigned to a processor. Ready processes are waiting to have the processor allocated to them by the operating system so that they can run. Process may come into this state after **Start** state or while running it by but interrupted by the scheduler to assign CPU to some other process.
- 3 **Running**
Once the process has been assigned to a processor by the OS scheduler, the process state is set to running and the processor executes its instructions.
- 4 **Waiting**
Process moves into the waiting state if it needs to wait for a resource, such as waiting for user input, or waiting for a file to become available.
- 5 **Terminated or Exit**
Once the process finishes its execution, or it is terminated by the operating system, it is moved to the terminated state where it waits to be removed from main memory.

Diagram of Process State



Process Control Block (PCB)

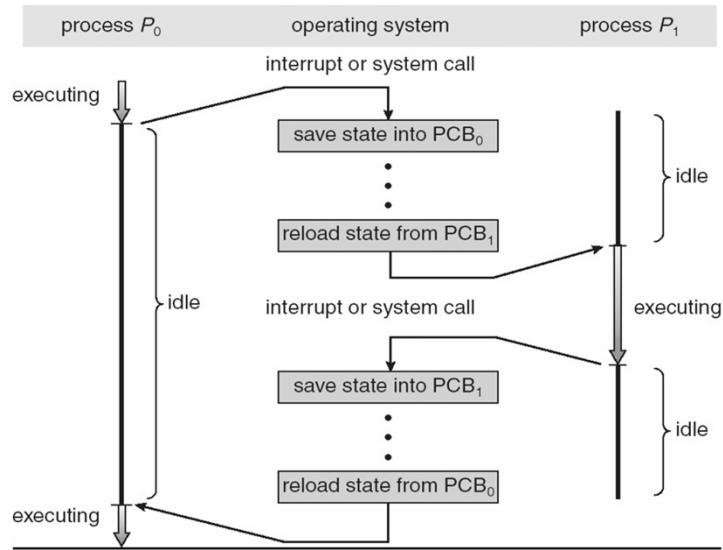
Information associated with each process

(also called **task control block**)

- ▶ Process state – running, waiting, etc
- ▶ Program counter – location of instruction to next execute
- ▶ CPU registers – contents of all process-centric registers
- ▶ CPU scheduling information- priorities, scheduling queue pointers
- ▶ Memory-management information – memory allocated to the process
- ▶ Accounting information – CPU used, clock time elapsed since start, time limits
- ▶ I/O status information – I/O devices allocated to process, list of open files

process state
process number
program counter
registers
memory limits
list of open files
...

CPU Switch From Process to Process



Context Switch

A **context switch** occurs when a computer's CPU **switches from one process or thread to a different process or thread**.

- Context switching allows for one CPU to **handle** numerous processes or threads **without** the **need** for additional **processors**.

- A context switch is the mechanism to **store and restore** the **state or context** of a CPU in **Process Control block** so that a process execution can be resumed from the same point at a later time.

- Any operating system that allows for multitasking relies heavily on the use of context switching to **allow different processes to run at the same time**.

Context switch

Typically, there are three situations that a context switch is necessary, as shown below.

- Multitasking** – When the CPU needs to switch processes in and out of memory, so that more than one process can be running.

- Kernel/User Switch** – When switching between user mode to kernel mode, it may be used (but isn't always necessary).

- Interrupts** – When the CPU is interrupted to return data from a disk read.

context switch

The steps in a full process switch are:

- 1.**Save the context** of the processor, including program counter and other registers.

- 2.**Update the process control block** of the process that is currently in the Running state. This includes changing the state of the process to one of the other states (Ready; Blocked; Ready/Suspend; or Exit). Other relevant fields must also be updated, including the reason for leaving the Running state and accounting information.

- 3.**Move the process control block** of this process to the **appropriate queue** (Ready; Blocked on Event *i* ; Ready/Suspend).

- 4.**Select another process** for execution.

- 5.Update the process control block of the process selected. This includes changing the state of this process to Running.

- 6.Update memory management data structures. This may be required, depending on how address translation is managed.

- 7.**Restore the context** of the processor to that which existed at the time the selected process was last switched out of the Running state, by loading in the previous values of the program counter and other registers.


Threads

- ▶ So far, process has a single thread of execution
- ▶ Consider having multiple program counters per process
 - ▶ Multiple locations can execute at once
 - ▶ Multiple threads of control -> **threads**
- ▶ Must then have storage for thread details, multiple program counters in PCB

Process Scheduling

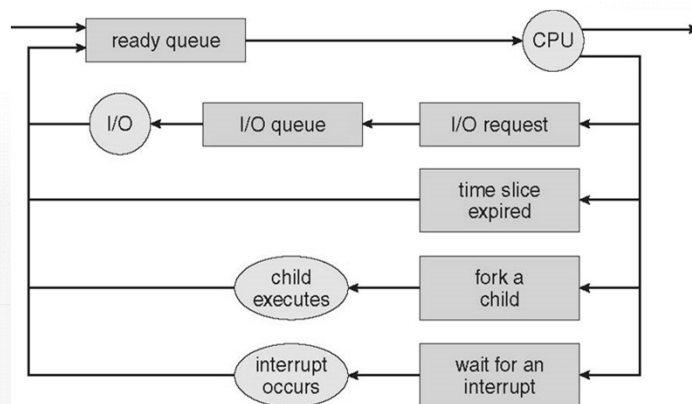
- ▶ Maximize CPU use, quickly switch processes onto CPU for time sharing
- ▶ **Process scheduler** selects among available processes for next execution on CPU
- ▶ Maintains **scheduling queues** of processes
 - ▶ **Job queue** – set of all processes in the system
 - ▶ **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - ▶ **Device queues** – set of processes waiting for an I/O device
 - ▶ Processes migrate among the various queues

Ready Queue And Various I/O Device Queues



Representation of Process Scheduling

- **Queueing diagram** represents queues, resources, flows



Process Scheduling

A new process is initially put in the ready queue. It waits there until it is selected for execution, or **dispatched**. Once the process is allocated the CPU and is executing, one of several events could occur:

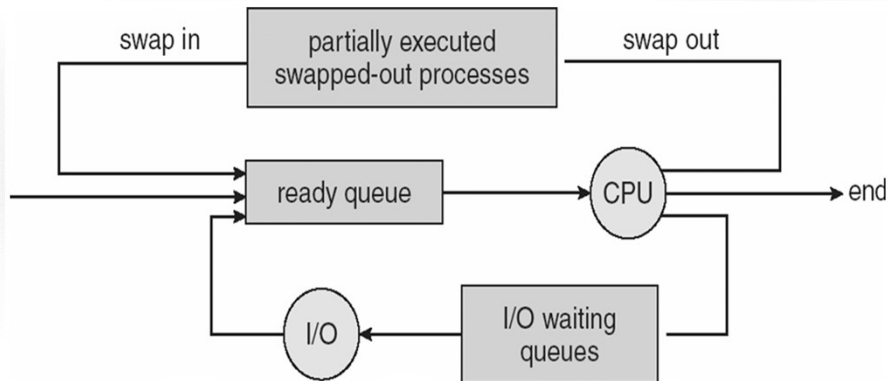
- The process could issue an I/O request and then be placed in an I/O queue.
- The process could create a new child process and wait for the child's termination.
- The process could be removed forcibly from the CPU, as a result of an interrupt, and be put back in the ready queue.

Schedulers

- ▶ **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
 - ▶ Sometimes the only scheduler in a system
 - ▶ Short-term scheduler is invoked frequently (milliseconds) ⇒ (must be fast)
- ▶ **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
 - ▶ Long-term scheduler is invoked infrequently (seconds, minutes) ⇒ (may be slow)
 - ▶ The long-term scheduler controls the **degree of multiprogramming**
- ▶ Processes can be described as either:
 - ▶ **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - ▶ **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- ▶ Long-term scheduler strives for good **process mix**

Addition of Medium Term Scheduling

- **Medium-term scheduler** can be added if degree of multiple programming needs to decrease
 - Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**



Comparison among Scheduler

S. N.	Long-Term Scheduler	Short-Term Scheduler	Medium-Term Scheduler
1	It is a job scheduler	It is a CPU scheduler	It is a process swapping scheduler.
2	Speed is lesser than short term scheduler	Speed is fastest among other two	Speed is in between both short and long term scheduler.
3	It controls the degree of multiprogramming	It provides lesser control over degree of multiprogramming	It reduces the degree of multiprogramming.
4	It is almost absent or minimal in time sharing system	It is also minimal in time sharing system	It is a part of Time sharing systems.
5	It selects processes from pool and loads them into memory for execution	It selects those processes which are ready to execute	It can re-introduce the process into memory and execution can be continued.

Multitasking in Mobile Systems

- ▶ Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended
- ▶ Due to screen real estate, user interface limits iOS provides for a
 - ▶ Single **foreground** process- controlled via user interface
 - ▶ Multiple **background** processes- in memory, running, but not on the display, and with limits
 - ▶ Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- ▶ Android runs foreground and background, with fewer limits
 - ▶ Background process uses a **service** to perform tasks
 - ▶ Service can keep running even if background process is suspended
 - ▶ Service has no user interface, small memory use

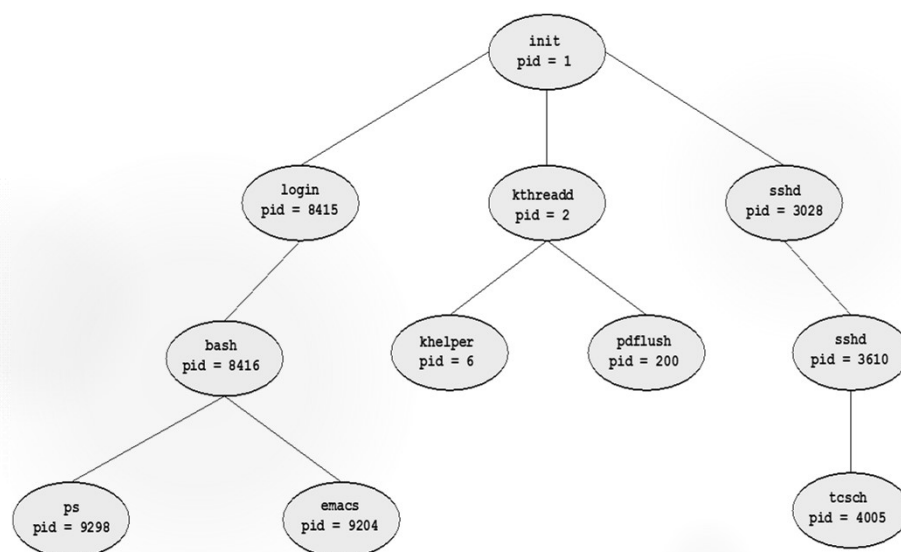
Operations on Processes

- ▶ System must provide mechanisms for:
 - ▶ process creation,
 - ▶ process termination,
 - ▶ and so on as detailed next

Process Creation


- ▶ **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- ▶ Generally, process identified and managed via a **process identifier (pid)**
- ▶ Resource sharing options
 - ▶ Parent and children share all resources
 - ▶ Children share subset of parent's resources
 - ▶ Parent and child share no resources
- ▶ Execution options
 - ▶ Parent and children execute concurrently
 - ▶ Parent waits until children terminate

A Tree of Processes in Linux



Process Creation (Cont.)

- ▶ Address space
 - ▶ Child duplicate of parent
 - ▶ Child has a program loaded into it
- ▶ UNIX examples
 - ▶ **fork()** system call creates new process
 - ▶ **exec()** system call used after a **fork()** to replace the process' memory space with a new program



C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Process Termination

- ▶ Process executes last statement and then asks the operating system to delete it using the **exit()** system call.
 - ▶ Returns status data from child to parent (via **wait()**)
 - ▶ Process' resources are deallocated by operating system
- ▶ Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
 - ▶ Child has exceeded allocated resources
 - ▶ Task assigned to child is no longer required
 - ▶ The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Process Termination

- ▶ Some operating systems do not allow child to exists if its parent has terminated. If a process terminates, then all its children must also be terminated.
 - ▶ **cascading termination.** All children, grandchildren, etc. are terminated.
 - ▶ The termination is initiated by the operating system.
- ▶ The parent process may wait for termination of a child process by using the **wait()** system call. The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```
- ▶ If no parent waiting (did not invoke **wait()**) process is a **zombie**
- ▶ If parent terminated without invoking **wait**, process is an **orphan**

Multiprocess Architecture – Chrome Browser

- ▶ Many web browsers ran as single process (some still do)
 - ▶ If one web site causes trouble, entire browser can hang or crash
- ▶ Google Chrome Browser is multiprocess with 3 different types of processes:
 - ▶ **Browser** process manages user interface, disk and network I/O
 - ▶ **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
 - ▶ Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - ▶ **Plug-in** process for each type of plug-in

